

LLM Architecture Evolution

From the Initial Transformer to microGPT, GPT-1, GPT-2, and GPT-3

Prepared with ChatGPT

July 25, 2026

Roadmap

- 1 The Initial Transformer (2017)
- 2 microGPT
- 3 GPT-1
- 4 GPT-2
- 5 GPT-3
- 6 Evolution Summary
- 7 References


Core Idea of the Initial Transformer

Main contribution

The paper *Attention Is All You Need* introduced the **Transformer** [Vaswani et al., 2017]: a sequence model built from self-attention + feed-forward layers + residual connections + normalization.

- Encoder–decoder architecture for sequence-to-sequence tasks such as translation.
- No recurrence and no convolution in the main sequence-processing path.
- Positional encoding is added because attention alone does not encode order.
- Decoder uses masked self-attention plus cross-attention to the encoder outputs.

Initial Transformer Architecture




What it is

microGPT is a tiny, dependency-free Python implementation of a decoder-only GPT. It is designed to expose the complete training and inference algorithm rather than maximize efficiency.

- Character-level language model on names.
- One Transformer layer in the example code: $L = 1$.
- Model width $d = 16$, number of heads $H = 4$, head size $d_h = 4$.
- Uses **RMSNorm** and a two-layer MLP $W_2 \text{ReLU}(W_1 x)$.
- Uses a learned token embedding E_{tok} and learned positional embedding E_{pos} .

microGPT Architecture



GPT-1: Generative Pre-Training

- Introduced in *Improving Language Understanding by Generative Pre-Training* (2018) [Radford et al., 2018].
- Decoder-only Transformer adapted from the original Transformer by discarding the encoder and cross-attention.
- Pretraining objective:


$$P(x_1, \dots, x_T) = \prod_{t=1}^T P(x_t \mid x_{<t}).$$

- Fine-tuning attaches task-specific heads after generative pretraining.

Typical GPT-1 configuration

Layers	12
Heads	12
Hidden size	768
Context	512

GPT-1 Architecture



GPT-2: Scaling the Decoder-Only GPT

Main message


GPT-2 kept the same basic decoder-only architecture as GPT-1, but scaled model size, dataset size, and context handling enough to demonstrate strong zero-shot and multitask behavior [Radford et al., 2019].

- Same core pattern: token embeddings → stack of masked self-attention blocks → LM head.
- Larger context length: 1024 tokens.
- Largest published model in the paper: about 1.5B parameters.
- Emphasis shifted from “pretrain then fine-tune” toward “pretrain and prompt/use directly.”

GPT-2 as an Enlarged GPT-1

GPT-2 small

12 layers, 12 heads,
768 hidden



GPT-3: Scaling to Few-Shot Learning

Main message


GPT-3 again retained the decoder-only Transformer pattern, but scaled it drastically [Brown et al., 2020]. The headline result was strong **in-context learning**: the model can perform new tasks from prompts alone, with zero, one, or a few demonstrations.

- Decoder-only Transformer with causal self-attention.
- Context length in the original paper: 2048 tokens.
- Largest model: 175B parameters.
- Major conceptual shift: prompting becomes a central interface to behavior.

GPT-3 Architecture: Same Skeleton, Much Larger Scale

Interpretation

Architecturally close to GPT-2, but large enough for few-shot and in-context learning to become prominent behaviors.



GPT

96 la
hidd
96 a
cont
175E

Architecture Evolution at a Glance

Model	Core architecture	Main objective	Characteristic scale	Main conceptual shift
Initial Transformer [Vaswani et al., 2017]	Encoder-decoder	Seq2seq prediction	6 encoder + 6 decoder layers in the paper	Attention replaces recurrence
microGPT	Tiny decoder-only GPT	Character-level next-token prediction	1 layer, 16 hidden, 4 heads	Educational minimal implementation
GPT-1 [Radford et al., 2018]	Decoder-only Transformer	Generative pretraining + fine-tuning	12 layers, 768 hidden	Transfer via generative pretraining
GPT-2 [Radford et al., 2019]	Larger decoder-only Transformer	Large-scale next-token prediction	up to 1.5B params	Zero-shot behavior emerges
GPT-3 [Brown et al., 2020]	Much larger decoder-only Transformer	Massive next-token prediction	175B params	Few-shot / in-context learning

The Main Structural Trend

From the initial Transformer to GPT-3

encoder-decoder Transformer $\implies$ decoder-only GPT family $\implies$ sam

- The original Transformer introduced the attention-based architecture.
- GPT keeps only the causal decoder side.
- microGPT strips that decoder-only architecture down to its conceptual essentials.
- GPT-1, GPT-2, and GPT-3 mostly differ by **scale**, training data, and the behavioral consequences of scale.

References I



Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin.

Attention Is All You Need.

arXiv:1706.03762, 2017.

<https://arxiv.org/abs/1706.03762>



Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever.

Improving Language Understanding by Generative Pre-Training.

OpenAI technical report, 2018.

https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf



Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever.

Language Models are Unsupervised Multitask Learners.

OpenAI technical report, 2019.

https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf

References II



Tom B. Brown et al.

Language Models are Few-Shot Learners.

arXiv:2005.14165, 2020.

<https://arxiv.org/abs/2005.14165>



Andrej Karpathy.

Minimal GPT-style educational code used as the microGPT reference in these slides.

User-provided code snippet in this study session.

LLM Architecture Evolution

From the initial Transformer to GPT-3

Questions?